

GETTING STARTED WITH THE ASSIGNMENT IN Python

We start with reading the credit scoring data from the lectures in Python.

The data can be found [here](#).

Suppose you saved the data to a file "credit.txt" in the directory "dm" on the C drive. To read it into Python type (">>>" denotes the prompt):

```
>>> import numpy as np
>>> credit_data = np.genfromtxt('C:/dm/credit.txt', delimiter=',', skip_header=True)
```

To display its value, just type its name at the command line:

```
>>> credit_data
array([[22., 0., 0., 28., 1., 0.],
       [46., 0., 1., 32., 0., 0.],
       [24., 1., 1., 24., 1., 0.],
       [25., 0., 0., 27., 1., 0.],
       [29., 1., 1., 32., 0., 0.],
       [45., 1., 1., 30., 0., 1.],
       [63., 1., 1., 58., 1., 1.],
       [36., 1., 0., 52., 1., 1.],
       [23., 0., 1., 40., 0., 1.],
       [50., 1., 1., 28., 0., 1.]])
```

"credit_data" is now a 2d NumPy array. Each rows represent a record and the columns represent the data attributes.

Select the first row of credit_data:

```
>>> credit_data[0]
array([22., 0., 0., 28., 1., 0.])
```

Select the fourth column of credit_data:

```
>>> credit_data[:,3]
array([28., 32., 24., 27., 32., 30., 58., 52., 40., 28.])
```

Select the element in row 4, column 0:

```
>>> credit_data[4,0]
29.0
```

Give the distinct values of income, sorted from low to high:

```
>>> np.sort(np.unique(credit_data[:,3]))
array([24., 27., 28., 30., 32., 40., 52., 58.])
```

Add all the entries of the sixth column:

```
>>> np.sum(credit_data[:,5])
5.0
```

Add the entries of each column of credit_data:

```
>>> credit_data.sum(axis=0)
array([363., 6., 7., 351., 5., 5.])
```

Add the entries of each row:

```
>>> credit_data.sum(axis=1)
array([ 51., 79., 51., 53., 63., 78., 125., 91., 65., 81.])
```

Select all rows where the first column is bigger than 27:

```
>>> credit_data[credit_data[:,0] > 27]
array([[46., 0., 1., 32., 0., 0.],
       [29., 1., 1., 32., 0., 0.],
       [45., 1., 1., 30., 0., 1.],
       [63., 1., 1., 58., 1., 1.],
       [36., 1., 0., 52., 1., 1.],
       [50., 1., 1., 28., 0., 1.]])
```

Construct a vector "x" with the numbers 2, 5, 10 in that order:

```
>>> x = np.array([2, 5, 10])
>>> x
array([ 2,  5, 10])
```

Construct a vector consisting of the numbers 0 through 9:

```
>>> np.arange(0, 10)
array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9])
```

Select the *row numbers* of the rows where the first column of `credit_data` is bigger than 27:

```
>>> np.arange(0, 10)[credit_data[:,0] > 27]
array([1, 4, 5, 6, 7, 9])
```

Draw a random sample of size 5 from the numbers 1 through 10 (without replacement):

```
>>> index = np.random.choice(np.arange(0, 10), size=5, replace=False)
>>> index
array([5, 7, 1, 3, 8])
```

Select the corresponding rows:

```
>>> train = credit_data[index,]
>>> train
array([[45.,  1.,  1., 30.,  0.,  1.],
       [36.,  1.,  0., 52.,  1.,  1.],
       [46.,  0.,  1., 32.,  0.,  0.],
       [25.,  0.,  0., 27.,  1.,  0.],
       [23.,  0.,  1., 40.,  0.,  1.]])
```

Select all rows with row number not in "index":
(This does not delete any rows from the original `credit_data`.)

```
>>> test = np.delete(credit_data, index, axis=0)
>>> test
array([[22.,  0.,  0., 28.,  1.,  0.],
       [24.,  1.,  1., 24.,  1.,  0.],
       [29.,  1.,  1., 32.,  0.,  0.],
       [63.,  1.,  1., 58.,  1.,  1.],
       [50.,  1.,  1., 28.,  0.,  1.]])
```

Consult the help page of the function "`np.random.choice`":

```
>>> help(np.random.choice)
```

Practice exercise 1

Assume we have a classification problem with only 2 classes that are labeled 0 and 1 respectively. Write a function that computes the impurity of a vector (of arbitrary length) of class labels. Use the gini-index as impurity measure. Do not use a loop structure in your function, this is not necessary.

Example:

```
>>> array=np.array([1,0,1,1,1,0,0,1,1,0,1])
>>> array
array([1,0,1,1,1,0,0,1,1,0,1])

>>> impurity(array)
0.23140495867768596
```

Practice exercise 2

Write a function "`bestsplit(x,y)`" that computes the best split value on a numeric attribute `x`. Here `x` is a vector of numeric values, and `y` is the vector of class labels (assume there are only two classes, coded as 0 and 1). `x` and `y` must be of the same length: `y[i]` is the class label of the `i`-th observation, and `x[i]` is the corresponding value of attribute `x`. Only consider splits of type "`x <= c`" where "`c`" is the average of two consecutive values of `x` in the sorted order. So one child contains all elements with "`x <= c`" and the other child contains all elements with "`x > c`". The best split is the split that achieves the highest impurity reduction.

Example (best split on income):

```
>>> bestsplit(credit_data[:,3],credit_data[:,5])
```

36

Hint: Clever use of "subscripting" (selecting elements of vectors and matrices) is important. For example, `y[x > 29]` produces a vector with all elements of `y` whose corresponding `x`-element (that is the element of `x` with the same index) is bigger than 29. More formally: `y[x > 29] = {y[i]: x[i] > 29}`. The result is a vector, not a set, i.e. duplicate values may occur. Just try it!

Hint: Example of how to determine candidate split points

```
>>> income_sorted = np.sort(np.unique(credit_data[:,3]))
>>> income_sorted
array([24, 27, 28, 30, 32, 40, 52, 58])

>>> income_splitpoints = (income_sorted[0:7]+income_sorted[1:8])/2
>>> income_splitpoints
array([25.5, 27.5, 29. , 31. , 36. , 46. , 55. ])
```

Note: use the "brute force" approach, i.e. don't implement the "segment borders" algorithm.